

```
In [1]: %matplotlib inline
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import r2_score
import time
import statsmodels.api as sm
from sklearn.linear_model import LinearRegression
```

Univariate Linear Regression

```
In [2]: data=pd.read_csv("Linear Regression/ex1data1.txt", header=None)
data.head()
```

Out[2]:

	0	1
0	6.1101	17.5920
1	5.5277	9.1302
2	8.5186	13.6620
3	7.0032	11.8540
4	5.8598	6.8233

```
In [3]: data.describe()
```

Out[3]:

	0	1
count	97.000000	97.000000
mean	8.159800	5.839135
std	3.869884	5.510262
min	5.026900	-2.680700
25%	5.707700	1.986900
50%	6.589400	4.562300
75%	8.578100	7.046700
max	22.203000	24.147000

```
In [4]: data.columns = ['Population', 'Profit'] #assigning column names
```

Profit Vs Population Graph

```
In [5]: plt.scatter(data['Population'], data['Profit'])
plt.xticks(np.arange(5,30,step=5))
plt.yticks(np.arange(-5,30,step=5))
plt.xlabel('Population (in 10,000s)')
plt.ylabel('Profit (in 10,000$)')
plt.title('Profit vs Population')
```

```
Out[5]: Text(0.5, 1.0, 'Profit vs Population')
```



Cost Function $J(\theta)$

```
In [6]: def computeCost(X,y,theta):
        """
        Take in a numpy array X,y,theta and get cost function using theta as parameter in a linear regression model
        """
        m = len(y)
        predictions = X.dot(theta)
        square_err = (predictions - y)**2

        return 1/(2*m)*np.sum(square_err)
```

Assigning X, y, θ and Reshaping the data

```
In [7]: data['x0'] = 1
data_val = data.values
m = len(data_val[:-1])
X = data[['x0', 'Population']].iloc[:-1].values
y = data['Profit'][:-1].values.reshape(m,1)
theta = np.zeros((2,1))

m, X.shape, y.shape, theta.shape
```

```
Out[7]: (96, (96, 2), (96, 1), (2, 1))
```

$$h(\theta) = x_0\theta_0 + x_1\theta_1 \dots (x_0 = 1)$$

```
In [8]: computeCost(X,y,theta)
```

```
Out[8]: 32.40484177877031
```

Gradient Descent

```
In [9]: def gradientDescent(X,y,theta,alpha,num_iters):
        """
        Take numpy array for X,y,theta and update theta for every iteration of gra
        dient steps
        Return theta and the list of cost of theta during each iteration
        """

        m = len(y)
        J_history = []
        for i in range(num_iters):
            predictions = X.dot(theta)
            error = np.dot(X.transpose(), (predictions-y))
            descent = alpha * 1/m * error
            theta -= descent
            J_history.append(computeCost(X,y,theta))

        return theta, J_history
```

```
In [10]: theta, J_history = gradientDescent(X,y,theta,0.001, 2000)
```


```
In [11]: print(f"h(x) = {str(round(theta[0,0],2))} + {str(round(theta[1,0],2))}x1")

h(x) = -1.11 + 0.92x1
```

```

In [12]: from mpl_toolkits.mplot3d import Axes3D
#Generating values for theta0, theta1 and the resulting cost value
theta0_vals=np.linspace(-10,10,100)
theta1_vals=np.linspace(-1,4,100)
J_vals=np.zeros((len(theta0_vals),len(theta1_vals)))
for i in range(len(theta0_vals)):
    for j in range(len(theta1_vals)):
        t=np.array([theta0_vals[i],theta1_vals[j]])
        J_vals[i,j]=computeCost(X,y,t)
#Generating the surface plot
fig = plt.figure()
ax = fig.add_subplot(111, projection='3d')
surf=ax.plot_surface(theta0_vals,theta1_vals,J_vals,cmap="coolwarm")
fig.colorbar(surf, shrink=0.5, aspect=5)
ax.set_xlabel("$\Theta_0$")
ax.set_ylabel("$\Theta_1$")
ax.set_zlabel("$J(\Theta)$")
#rotate for better angle
ax.view_init(30,120)


```



Cost Graph


```
In [13]: plt.plot(J_history)
plt.xlabel("Iteration")
plt.ylabel("$J(\Theta)$")
plt.title("Cost function using Gradient Descent")
```

Out[13]: Text(0.5, 1.0, 'Cost function using Gradient Descent')



```
In [14]: plt.scatter(data['Population'], data['Profit'])
x_value = [x for x in range(25)]
y_value = [x*theta[1] + theta[0] for x in x_value]
plt.plot(x_value, y_value, color = 'r')
plt.xticks(np.arange(5,30,step=5))
plt.yticks(np.arange(-5,30,step=5))
plt.xlabel('Population (in 10,000s)')
plt.ylabel('Profit (in 10,000$)')
plt.title('Profit vs Population')
```

Out[14]: Text(0.5, 1.0, 'Profit vs Population')



Function to predict

```
In [15]: def predict(x,theta):
        """
        Takes in numpy array x and theta and returns predicted value of y
        """
        predictions = np.dot(theta.transpose(),x)
        return predictions[0]
```

```
In [16]: data.tail(1)
```

Out[16]:

	Population	Profit	x0
96	5.4369	0.61705	1

```
In [17]: predict1 = predict(data[['x0', 'Population']].iloc[-1].values, theta)*10000
        print(f'For a population of 6170 the predicted profit is ${predict1}')
```

For a population of 6170 the predicted profit is \$38686.246103378166

TO DO: Configure code for

Multivariate Linear Regression

```
In [18]: data2=pd.read_csv("Linear Regression/ex1data2.txt",header=None)
        data2.head()
```

Out[18]:

	0	1	2
0	2104	3	399900
1	1600	3	329900
2	2400	3	369000
3	1416	2	232000
4	3000	4	539900

```
In [19]: data2.columns=['Size of house', 'No. of bedrooms', 'Price']  
data2.head()
```

Out[19]:

	Size of house	No. of bedrooms	Price
0	2104	3	399900
1	1600	3	329900
2	2400	3	369000
3	1416	2	232000
4	3000	4	539900

```
In [20]: data2.describe()
```

Out[20]:

	Size of house	No. of bedrooms	Price
count	47.000000	47.000000	47.000000
mean	2000.680851	3.170213	340412.659574
std	794.702354	0.760982	125039.899586
min	852.000000	1.000000	169900.000000
25%	1432.000000	3.000000	249900.000000
50%	1888.000000	3.000000	299900.000000
75%	2269.000000	4.000000	384450.000000
max	4478.000000	5.000000	699900.000000

Price Vs Size of house & Price Vs No. of bedrooms


```
In [21]: plt.scatter(data2['Size of house'], data2['Price'])
plt.xticks(np.arange(500, 5000, step=500))
plt.yticks(np.arange(100000, 800000, step=50000))
plt.xlabel('Size of house')
plt.ylabel('Price')
plt.title('Price vs Size of house')
```

Out[21]: Text(0.5, 1.0, 'Price vs Size of house')



```
In [22]: plt.scatter(data2['No. of bedrooms'], data2['Price'])
plt.xticks(np.arange(0, 6, step=1))
plt.yticks(np.arange(100000, 800000, step=50000))
plt.xlabel('No. of bedrooms')
plt.ylabel('Price')
plt.title('Price vs No. of bedrooms')
```

Out[22]: Text(0.5, 1.0, 'Price vs No. of bedrooms')



Since column "size of house" has very high values we need to normalize it

Scaling column "size of house"

```
In [23]: col = ['Size of house', 'Price']
scaler=MinMaxScaler()
data2[col]=pd.DataFrame(scaler.fit_transform(data2[col]), columns=data2 [col].
columns)
```

Splitting X and y from original dataset (also into train and test)

```
In [24]: y = np.array(data2['Price'][:-1])
X = np.array(data2.drop('Price', axis =1)[:-1])
X.shape, y.shape
```

```
Out[24]: ((46, 2), (46,))
```

We need y shape as (46,1) and add a column of 1's in X

Reshaping X,y

```
In [25]: y=y.reshape(y.shape[0],1)
X=np.c_[np.ones(X.shape[0]),X]
X.shape,y.shape
```

```
Out[25]: ((46, 3), (46, 1))
```

Configuring the functions (*costCompute, gradientDescent, predict*) for Multivariate

I am writing code for costCompute and gradientDescent as taught by Siba Sir.

```
In [26]: def computeCost(X,Y,theta):
        """
        Take in a numpy array X,y,theta and get cost function using theta as param
        eter in a linear regression model
        """
        m=Y.size
        h_theta=np.dot(X,theta) #Predictions
        error=h_theta-Y
        cost=(1/(2*m))*np.dot(error.T,error)
        return cost
```

```
In [27]: def gradientDescent(X,Y,theta,alpha,num_iters):
        """
        Take numpy array for X,y,theta and update theta for every iteration of gradient steps
        Return theta and the list of cost of theta during each iteration
        """
        m = Y.size
        past_cost = []
        past_theta=[theta]
        for i in range(num_iters):
            h_theta = np.dot(X,theta)
            error=h_theta-Y
            past_cost.append(computeCost(X,Y,theta))
            theta=theta-alpha*(1/m)*np.dot(X.T,error)
            past_theta.append(theta)
            if(past_theta[i]==past_theta[i+1]).all() : break
        return past_theta, past_cost,i
```

```
In [28]: def predict(X,theta):
        """
        Takes in numpy array x and theta and returns predicted value of y
        """
        predictions = np.dot(X,theta)
        return predictions
```

Assigning random values for θ and values for α & number of iterations we want

```
In [29]: np.random.seed(123)
        theta=np.random.rand(X.shape[1],1)
        alpha=0.05
        num_iters=50000
```

Performing Linear Regression and displaying results

```
In [30]: start=time.time()
        past_theta,past_cost,stop=gradientDescent(X,y,theta,alpha,num_iters)
        timetaken=time.time()-start
```

```
In [31]: best_theta=past_theta[-1]
        best_cost=past_cost[-1]
```

```
In [32]: print(f'Our model performed {stop} iterations out of {num_iters} iterations before converging (the previous  $\theta$  was equal to current  $\theta$ )')
```

Our model performed 21483 iterations out of 50000 iterations before converging (the previous θ was equal to current θ)

```
In [33]: print(f'Best Theta: \n {best_theta}')
         print(f'Best Cost: \n {best_cost}')
```

```
Best Theta:
[[ 0.07198606]
 [ 0.95458358]
 [-0.01672784]]
Best Cost:
[[0.00742914]]
```

```
In [34]: print(f'h(θ) = {str(round(best_theta[0,0],5))} + {str(round(best_theta[1,0],
5))}x1 + {str(round(best_theta[2,0],5))}x2')
         print(f'\nTime taken :{timetaken}')
         print(f'\nAccuracy: {round(r2_score(y,predict(X,best_theta)),4)*100}%')
```

$h(\theta) = 0.07199 + 0.95458x_1 + -0.01673x_2$

Time taken :0.501678466796875


Accuracy: 72.91%

Plotting Cost

```
In [35]: cost = np.asarray(past_cost)
         cost = cost.reshape((cost.shape[0],1))
         cost.shape
```

Out[35]: (21484, 1)

```
In [36]: plt.plot(cost)
         plt.show()
```



Predicting for test data

```
In [37]: data2.tail(1)
```

```
Out[37]:
```

	Size of house	No. of bedrooms	Price
46	0.096801	3	0.131321

```
In [38]: predict2=predict(X[-1], best_theta)*1000000
print(f'For house of size 968 sq.ft. the predicted price is :${predict2}')
```

For house of size 968 sq.ft. the predicted price is :\$[268335.49803214]

Comparing parameters

```
In [39]: print(f'Parameters from StatsModels -> {sm.OLS(y,X).fit().params}')
print(f'Parameters from ScikitLearn -> {LinearRegression().fit(X,y).coef_}')
print(f'Parameters from Gradient Descent -> {best_theta.reshape(3,)}')
```

Parameters from StatsModels -> [0.07198606 0.95458358 -0.01672784]
 Parameters from ScikitLearn -> [[0. 0.95458358 -0.01672784]]
 Parameters from Gradient Descent -> [0.07198606 0.95458358 -0.01672784]

Here we can see that the results match

Code for Linear Regression

if we want to put the splitting data into X,y and its reshaping into a function too

```
In [40]: def LinearReg_GD (data, pred_col, alpha, num_iters, dec = 5):
#Declaring X,y and theta
y1 = np.array(data[pred_col]).reshape((-1,1))
X1 = np.array(data.drop(pred_col, axis = 1))
X1 = np.c_[np.ones (X1.shape[0]), X1]
theta = np.zeros((X1.shape[1],1))
return gradientDescent(X1,y1,theta,alpha,num_iters)
```

```
In [41]: np.random.seed(123)
theta1=np.random.rand(X.shape[1],1)
alpha=0.05
num_iters=50000
past_theta1,past_cost1,stop1=LinearReg_GD(data2, 'Price', alpha,num_iters)
```

```
In [42]: best_theta1=past_theta1[-1]
best_cost1=past_cost1[-1]

print(f'Our model performed {stop1} iterations out of {num_iters} iterations before converging (the previous  $\theta$  was equal to current  $\theta$ )\n\n')

print(f'Best Theta: \n {best_theta1}')
print(f'Best Cost: \n {best_cost1}')
```

```
print(f'\n\nh( $\theta$ ) = {str(round(best_theta1[0,0],5))} + {str(round(best_theta1[1,0],5))}x1 + {str(round(best_theta1[2,0],5))}x2')
```

Our model performed 22195 iterations out of 50000 iterations before converging (the previous θ was equal to current θ)

Best Theta:
[[0.07227435]
[0.95241114]
[-0.01648683]]
Best Cost:
[[0.00727405]]

$h(\theta) = 0.07227 + 0.95241x_1 + -0.01649x_2$